


Variational Inference

Variational Inference



Variational Inference: Idea

we are generally interested in getting insights from observed data

inferring latent variables from observed data

a conditional probabilistic distribution

$$x \sim p(x) \quad z \sim p(z|x) \quad (\text{posterior}) \quad (\text{inference})$$

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)} = \frac{p(x|z)p(z)}{\int_z p(x, z) dz} \quad p(x) \text{ is generally intractable}$$

idea: use a family of easy distribution to approximate the hard distribution

(variational)




$$\begin{aligned} z &\sim p(z) \\ x &\sim p(x|z) \\ &(\text{generation}) \end{aligned}$$

Variational Inference: Example (1)

$$p(z|x) = \frac{p(x|z)p(z)}{p(x)}$$

bayes' rule



$$z \sim p(z) = \begin{cases} e^{-z}, & z \geq 0 \\ 0, & z < 0 \end{cases} = e^{-z} I(z \geq 0)$$

$$x \sim p(x|z) = N(x, \mu = z, \sigma = 1) = \frac{1}{\sqrt{2\pi}} e^{(-\frac{1}{2}(x-z)^2)}$$

$$p(x, z) = p(x|z)p(z) = \frac{1}{\sqrt{2\pi}} e^{(-\frac{1}{2}(x-z)^2)} e^{-z} I(z \geq 0) \quad (\text{joint distribution})$$

$$p(x) = \int_0^\infty p(x, z) dz = \int_0^\infty e^{-z} \frac{1}{\sqrt{2\pi}} e^{(-\frac{1}{2}(x-z)^2)} dz$$

thus the posterior has no closed-form solution

the integral has no closed-form solution

Latent Graphical Model Variational Inference: Idea Variational Inference: Example Setup Variational Inference: KL, ELBO & Example Take Away

Variational Inference: Example (2)

$$\begin{aligned} p(z|x) &\sim p(x, z) = p(x|z)p(z) = \frac{1}{\sqrt{2\pi}} e^{(-\frac{1}{2}(x-z)^2)} e^{-z} I(z \geq 0) \\ &= \frac{1}{\sqrt{2\pi}} e^{(-\frac{1}{2}z^2 + (x-1)z - \frac{1}{2}x^2)} I(z \geq 0) \\ &= \frac{1}{\sqrt{2\pi}} e^{(-\frac{1}{2}(z-(x-1))^2 + \frac{1}{2}(x-1)^2 - \frac{1}{2}x^2)} I(z \geq 0) \\ &= \frac{1}{\sqrt{2\pi}} e^{(\frac{1}{2}(x-1)^2 - \frac{1}{2}x^2)} e^{(-\frac{1}{2}(z-(x-1))^2)} I(z \geq 0) \\ &\sim \frac{1}{\sqrt{2\pi}} e^{(-\frac{1}{2}(z-(x-1))^2)} I(z \geq 0) \end{aligned}$$


posterior proportional to a gaussian curve for z larger or equal to zero

Latent Graphical Model Variational Inference: Idea Variational Inference: Example Setup Variational Inference: KL, ELBO & Example Take Away


Variational Inference: Example (3)



$$p(z|x) \sim \frac{1}{\sqrt{2\pi}} e^{-\frac{1}{2}(z-(x-1))^2} I(z \geq 0)$$



$x = 1.5$



$x = 0$

let's now use an “easier” distribution to approximate the intractable posterior

Latent Graphical Model Variational Inference: Idea Variational Inference: Example Setup Variational Inference: KL, ELBO & Example Take Away

Approximate Distribution: ELBO



$$q_{\theta}(z) \approx p(z|x)$$

$$\begin{aligned} D(q_{\theta}(z) \| p(z|x)) &= E_{z \sim q} \left[\log \frac{q_{\theta}(z)}{p(z|x)} \right] = E_{z \sim q} [\log q_{\theta}(z) - \log p(z|x)] \\ &= E_{z \sim q} [\log q_{\theta}(z) - \log \frac{p(z, x)}{p(x)}] \\ &= E_{z \sim q} [\log q_{\theta}(z) - \log p(z, x)] + \log p(x) \end{aligned}$$

evidence lower bound (elbo)

$$\log p(x) = E_{z \sim q} [\log p(z, x) - \log q_{\theta}(z)] + D(q_{\theta}(z) \| p(z|x))$$

$$\log p(x) \geq E_{z \sim q} [\log p(z, x) - \log q_{\theta}(z)] \equiv \mathcal{L}_q$$



objective: minimizing kl divergence \rightarrow maximizing elbo (variational optimization: optimizing over functions)

Latent Graphical Model Variational Inference: Idea Variational Inference: Example Setup Variational Inference: KL, ELBO & Example Take Away

Note: here (in standard variational inference), we know $p(z|x)$, which is the real distribution, because we know $p(x)$ and $p(z, x)$. But in VAE, which is introduced later, we don't know $p(z|x)$, so we use $q(z|x)$ to approximate it, which is called amortized variational inference. $q(z|x)$ is the conditional probability the encoder (the neural network) gives, which is based on a signal data piece. We have many pieces of data to derive different $q(z|x)$ and optimize it (called “amortized”).

$q(z|x)$ is a variational distribution, and is commonly a Gaussian parametrized by the encoder network outputs.



Take Away



an inference problem can be formulated as a **latent graphical model** where given the observed data we would like to infer the **posterior probability distribution** of associated **latent variables**, which can be usually interpreted as properties

the posterior is typically intractable due to typically the marginal requires an integral to compute

variational inference uses a simple, parameteric distribution to approximate the true posterior and turns the intractable problem into an optimization problem

kl divergence measures the difference between the approximate distribution and the true posterior

minimizing kl divergence is equivalent of maximizing evidence lower bound (elbo)

elbo is negative & plays a very important role in subsequent algorithm formulations

Latent Graphical Model Variational Inference: Idea Variational Inference: Example Setup Variational Inference: KL, ELBO & Example Take Away

Variational AutoEncoder

ELBO (1)

original form

$$\mathcal{L}_q = E_{z \sim q}[\log p(z, x) - \log q(z)]$$

$q(z) \approx p(z|x)$ approximate distribution

z latent variable x observed variable

variational autoencoder changes

approximation distribution $q_\phi(z|x)$ encoder (amortized variational inference)

likelihood distribution $p_\theta(x|z)$ decoder

$$\begin{aligned}\mathcal{L}_q &= E_{z \sim q}[\log p_\theta(x|z)p(z) - \log q_\phi(z|x)] \\ &= E_{z \sim q}[\log p_\theta(x|z) + \log p(z) - \log q_\phi(z|x)]\end{aligned}$$



intuition ELBO Objective Reparameterization Loss Function Visualizations Take Away

ELBO (2)

Press **Esc** to exit full screen


$$\begin{aligned}\mathcal{L}_q &= E_{z \sim q}[\log p_\theta(x|z) + \log p(z) - \log q_\phi(z|x)] \\ &= E_{z \sim q}[\log p_\theta(x|z)] - E_{z \sim q}[\log \frac{q_\phi(z|x)}{p(z)}] \\ &= E_{z \sim q}[\log p_\theta(x|z)] - D(q_\phi(z|x) \| p(z))\end{aligned}$$

$E_{z \sim q}[\log p_\theta(x|z)]$ reconstruction likelihood

$D(q_\phi(z|x) \| p(z))$ latent prior similarity
 $p(z) = N(0, I)$

maximizing reconstruction likelihood
while minimizing kl divergence

(approximating $\log p(x)$)



ref. 2

intuition ELBO Objective Reparameterization Loss Function Visualizations Take Away

Objective (1)

$$q_{\varphi}(z|x) = N(z; \mu_{\varphi}(x), \Sigma_{\varphi}(x))$$

$$p(z) = N(z; 0, I)$$

kl divergence of gaussian

$$D(N(\mu_0, \Sigma_0) \| N(\mu_1, \Sigma_1)) = \frac{1}{2} (\text{tr}(\Sigma_1^{-1} \Sigma_0) + (\mu_1 - \mu_0)^T \Sigma_1^{-1} (\mu_1 - \mu_0) - k + \log(\frac{\det \Sigma_1}{\det \Sigma_0}))$$

$$D(q_{\varphi}(z|x) \| p(z)) = \frac{1}{2} (\text{tr}(\Sigma_{\varphi}(x)) + \mu_{\varphi}(x)^T \mu_{\varphi}(x) - k - \log(\det \Sigma_{\varphi}(x))) \quad k \text{ is the dimension}$$

encoder objective

Objective (2)

$$p_{\theta}(x|z) = N(x; \mu_{\theta}(z), \sigma^2 I)$$

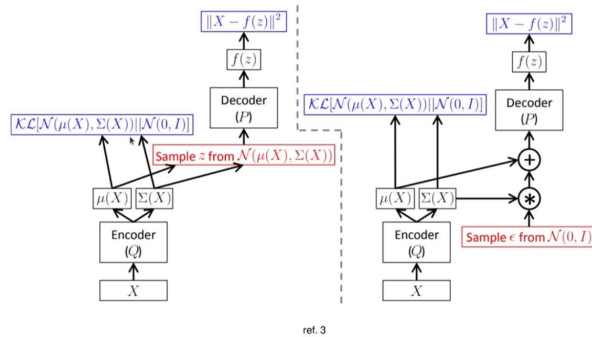
$$\log p_{\theta}(x|z) = -c \|x - \mu_{\theta}(z)\|^2 + d \quad \nabla_{\theta} \|x - \mu_{\theta}(z)\|^2 \quad \text{decoder objective}$$

but

$$E_{z \sim q_{\varphi}} [\log p_{\theta}(x|z)] \quad q_{\varphi}(z|x) = N(z; \mu_{\varphi}(x), \Sigma_{\varphi}(x))$$

cannot back propagate

Reparameterization



$$\begin{aligned} q_{\phi}(z|x) &= N(z; \mu_{\phi}(x), \Sigma_{\phi}(x)) \\ &= \mu_{\phi}(x) + \Sigma_{\phi}^{1/2}(x) N(0, I) \\ &= \mu_{\phi}(x) + \Sigma_{\phi}^{1/2}(x) \epsilon \end{aligned}$$

$$E_{\epsilon \sim N(0, I)} [\log p_{\theta}(x|z = \mu_{\phi}(x) + \Sigma_{\phi}^{1/2}(x) \epsilon)]$$



Loss Function

reconstruction loss

$$\mathcal{L}_1 = -\log p_{\theta}(x|z) \sim \|x - \mu_{\theta}(z)\|^2$$

latent prior similarity loss


$$\mathcal{L}_2 = D(q_{\phi}(z|x) \| p(z)) = \frac{1}{2} (\text{tr}(\Sigma_{\phi}(x)) + \mu_{\phi}(x)^T \mu_{\phi}(x) - k - \log(\det \Sigma_{\phi}(x)))$$

total loss

$$\mathcal{L} = \mathcal{L}_1 + \alpha \mathcal{L}_2$$

Latent Space Visualizations


up讲的很好!



VAE on MNIST

Intuition ELBO Objective Reparameterization Loss Function Visualizations Take Away

Image Generation Visualizations



ref. 1

VAE on MNIST

Intuition ELBO Objective Reparameterization Loss Function Visualizations Take Away

Take Away



variational autoencoder can be understood as an instantiation of variational inference with encoder-decoder based framework

latent space prior is chosen to be unit gaussian

variational autoencoder can be used to generate data by sampling latent variables from unit gaussian

loss function can be derived from elbo which constitutes reconstruction loss and latent prior similarity loss, which can be treated as a regularization term

reparameterization trick is used to enable proper backpropagation

latent space visualization shows the latent prior constraint makes the learned latent variable distribution relatively continuous and has somewhat of a topology in it